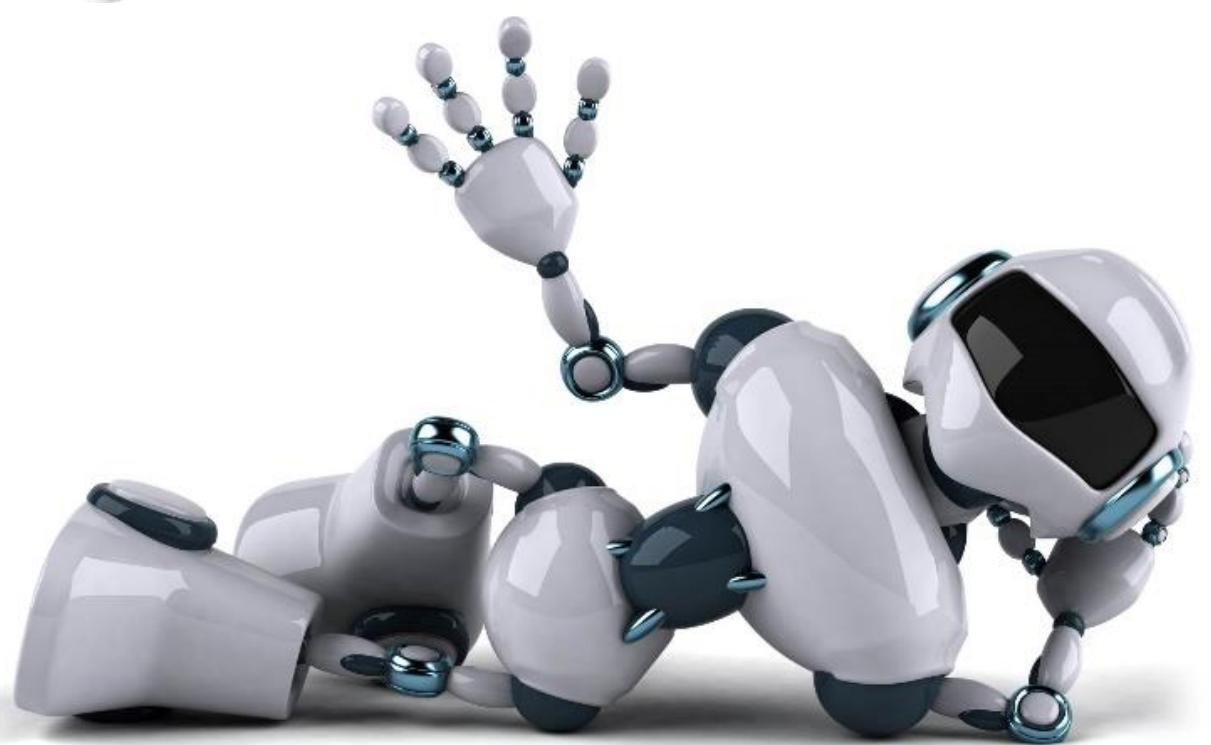


Intro to ML for Input Recognition

CHI Course Computational Interaction 2016

Otmar Hilliges

10 April 2017



Designing Interactive Systems

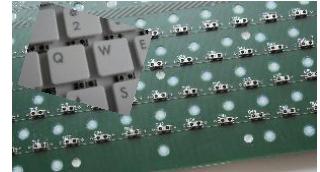
Interaction
Task



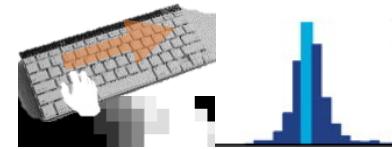
Sensing
Technology



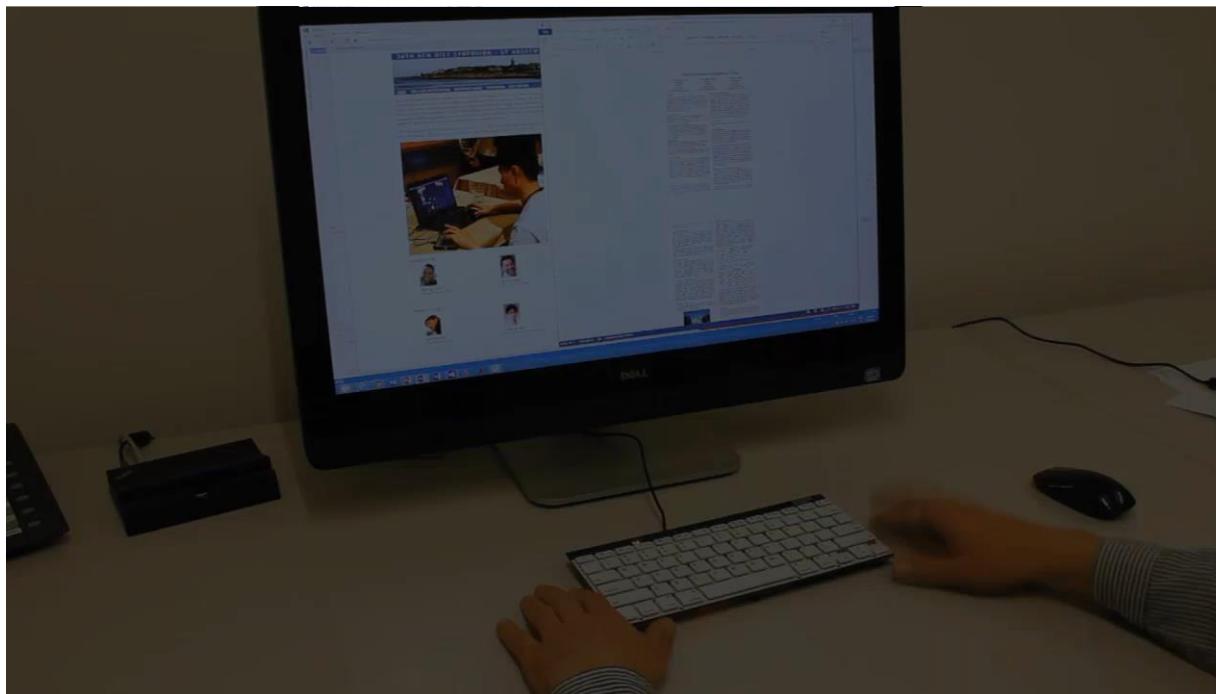
Sensor
Placement



Input
Recognition



System
Output



[Motion Sensing Keyboard. ACM CHI'14]

Today's talk – Part I

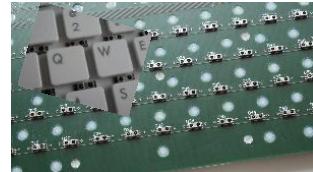
Interaction
Task



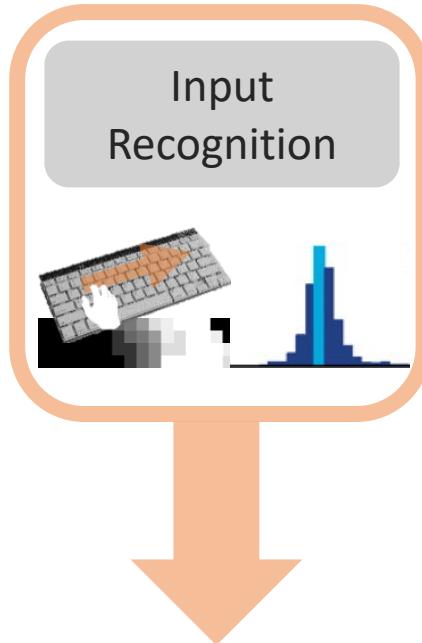
Sensing
Technology



Sensor
Configuration



Input
Recognition



System
Output



Today: Machine Learning for Gesture Recognition



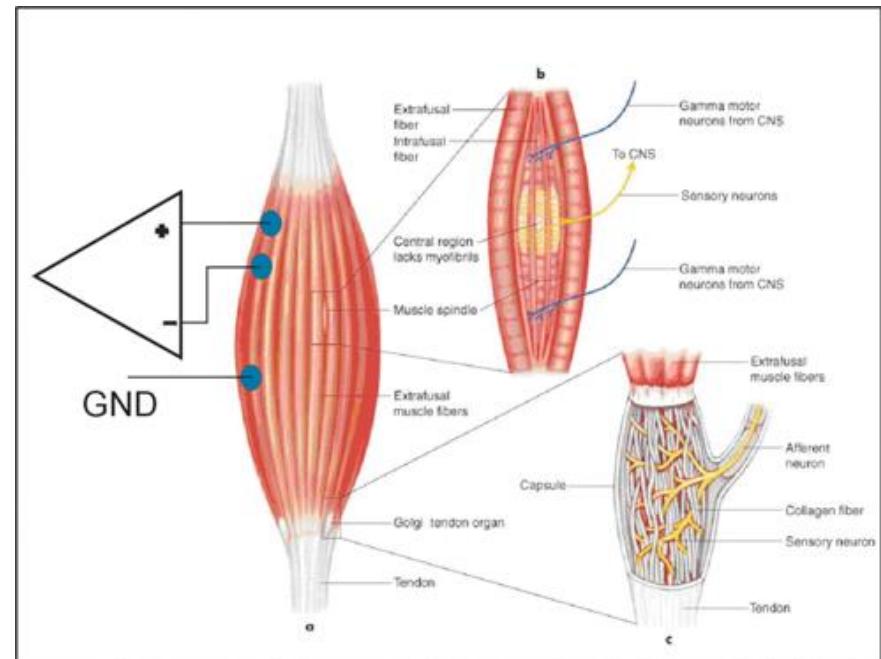
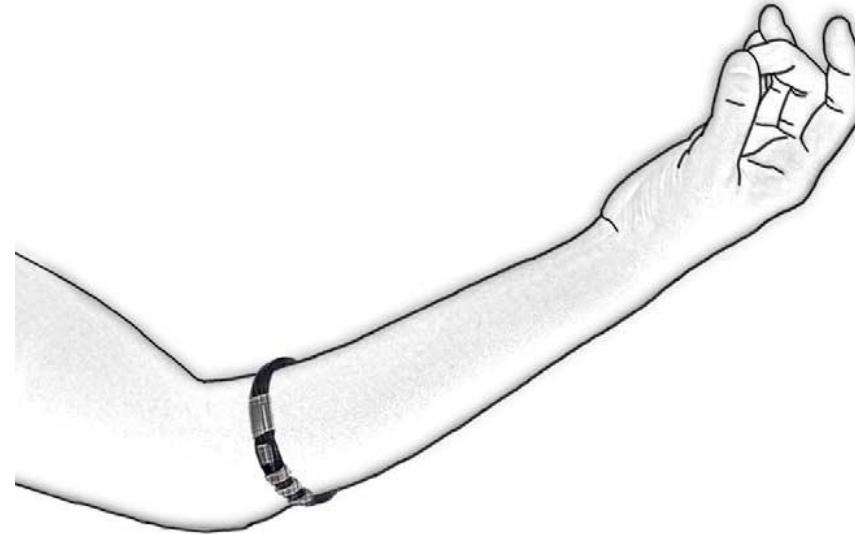
Muscle-Computer Interfaces

Hands free interaction – even when holding an object

EMG or Electromyography
(primarily used in medical therapy)

Action potential generated by muscle when signal arrives from motor neuron

Can be sensed non-invasively by placing electrodes on the skin



Video



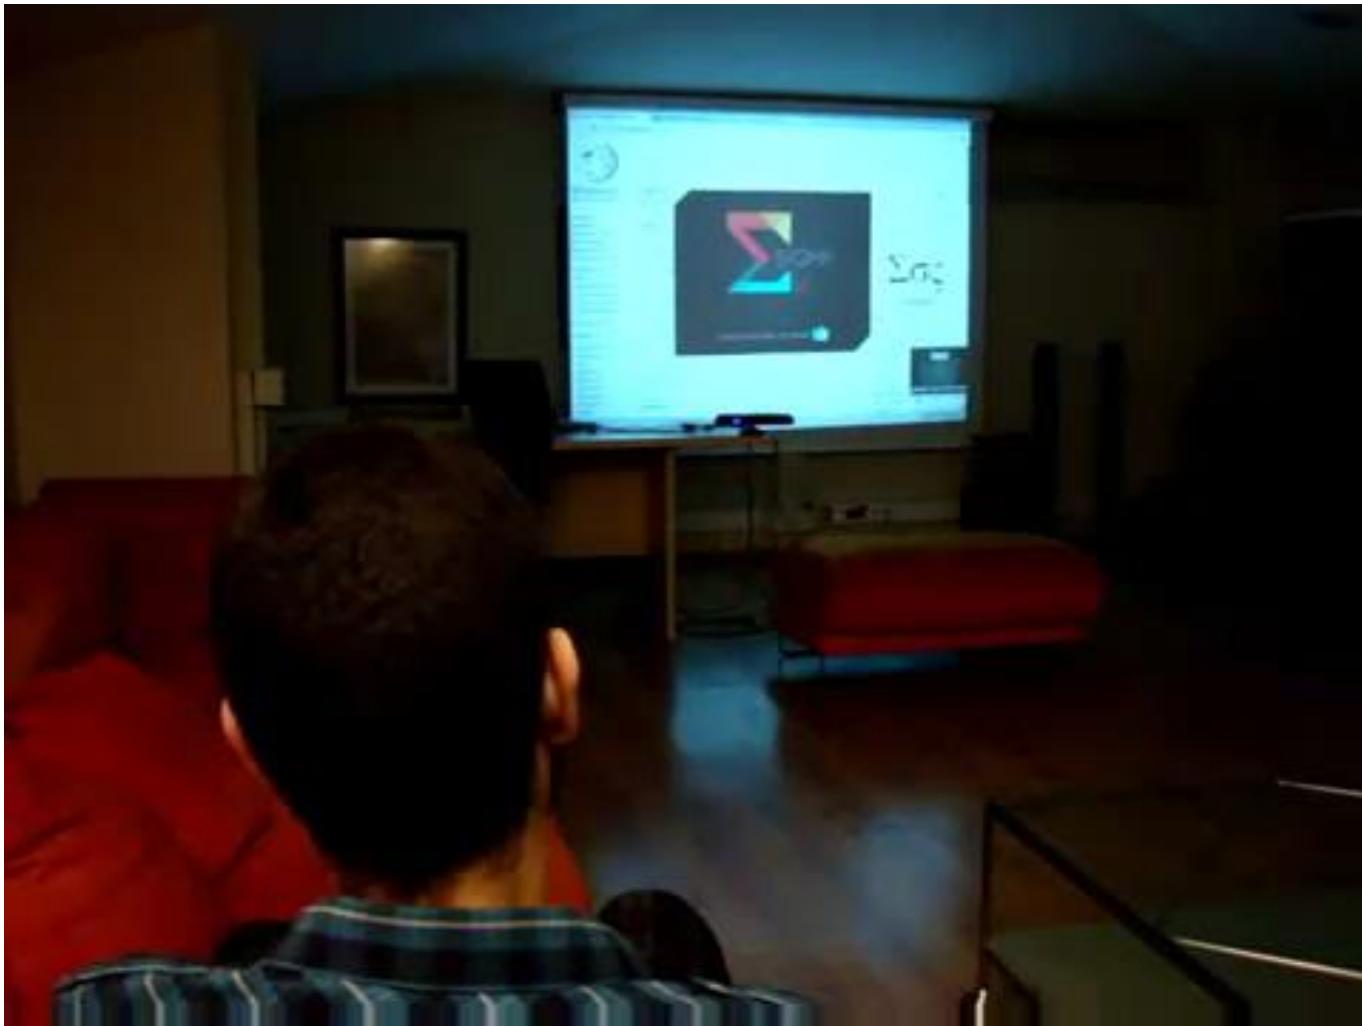
[Saponas et al., *Proceedings of ACM UIST '09*]

Touch & Activate: Video

Touch & Activate: Adding Interactivity to Existing Objects
using Active Acoustic Sensing

Makoto Ono, Buntarou shizuki, and Jiro Tanaka
University of Tsukuba

Gesture Recognition



<http://www.sigmaro.com/portfolio/sigmanil-framework/>

Take Home Message

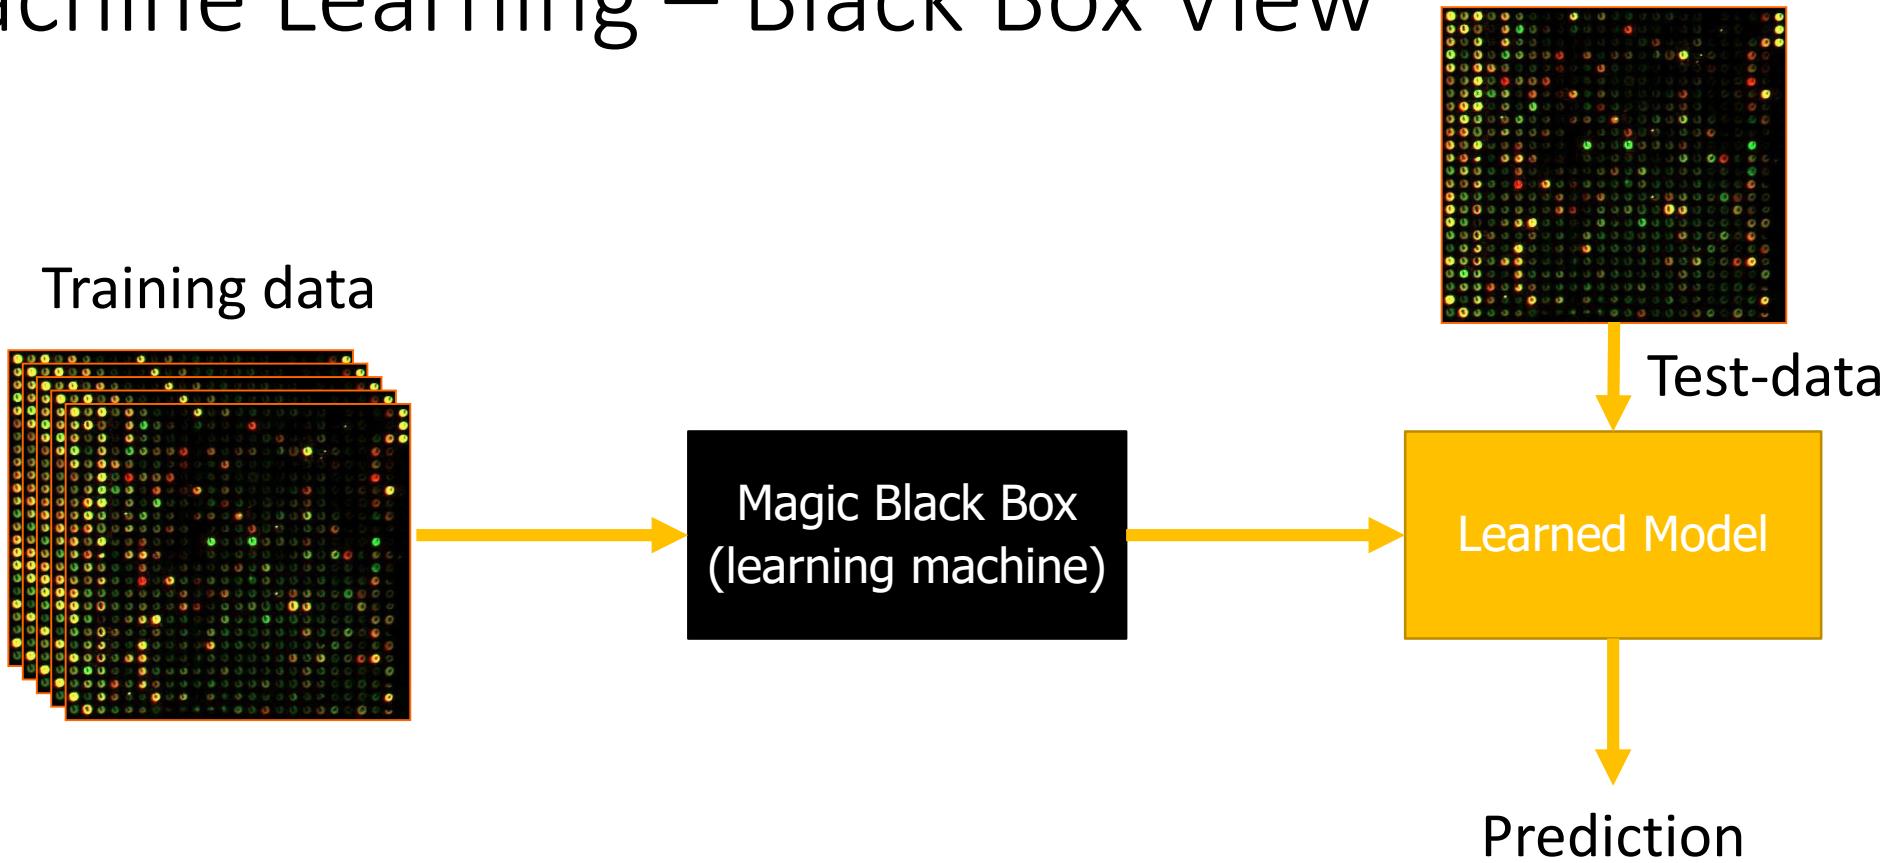
Humans use a variety of channels to interact & communicate

- Human behavior is highly complex
- Human behavior is highly ambiguous
- To create we intuitive interfaces we want to use:
 - Small, low-effort natural motion
 - Avoid heavy user instrumentation

In consequence,

- Recognizing human input is challenging
- Most heuristic / traditional approaches are not powerful enough to model the variance in (gesture) execution
- Many different sensor modalities, often noisy data

Machine Learning – Black Box View



Training data:

- Positive *and* negative examples. Data has to be labelled (ground truth).

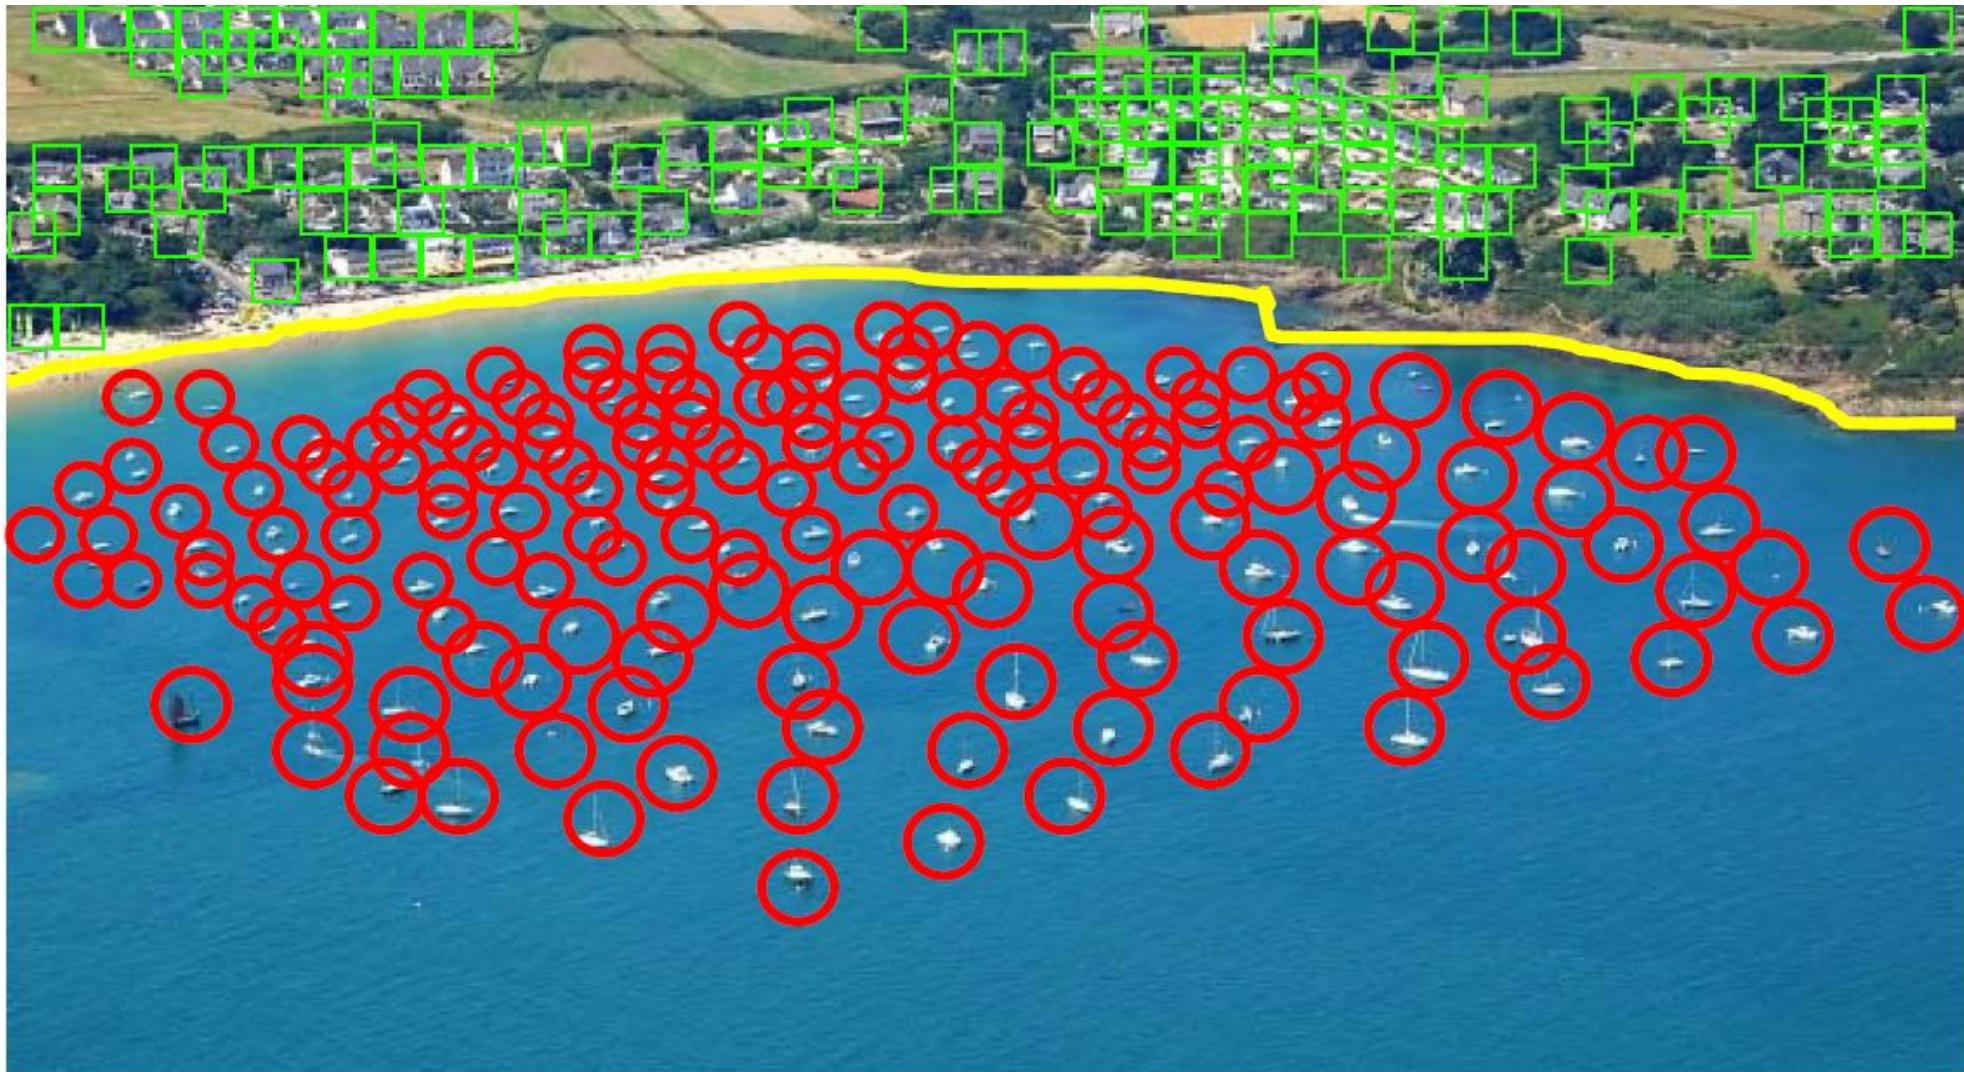
Model:

- The model can distinguish between classes or estimate real valued outcome.

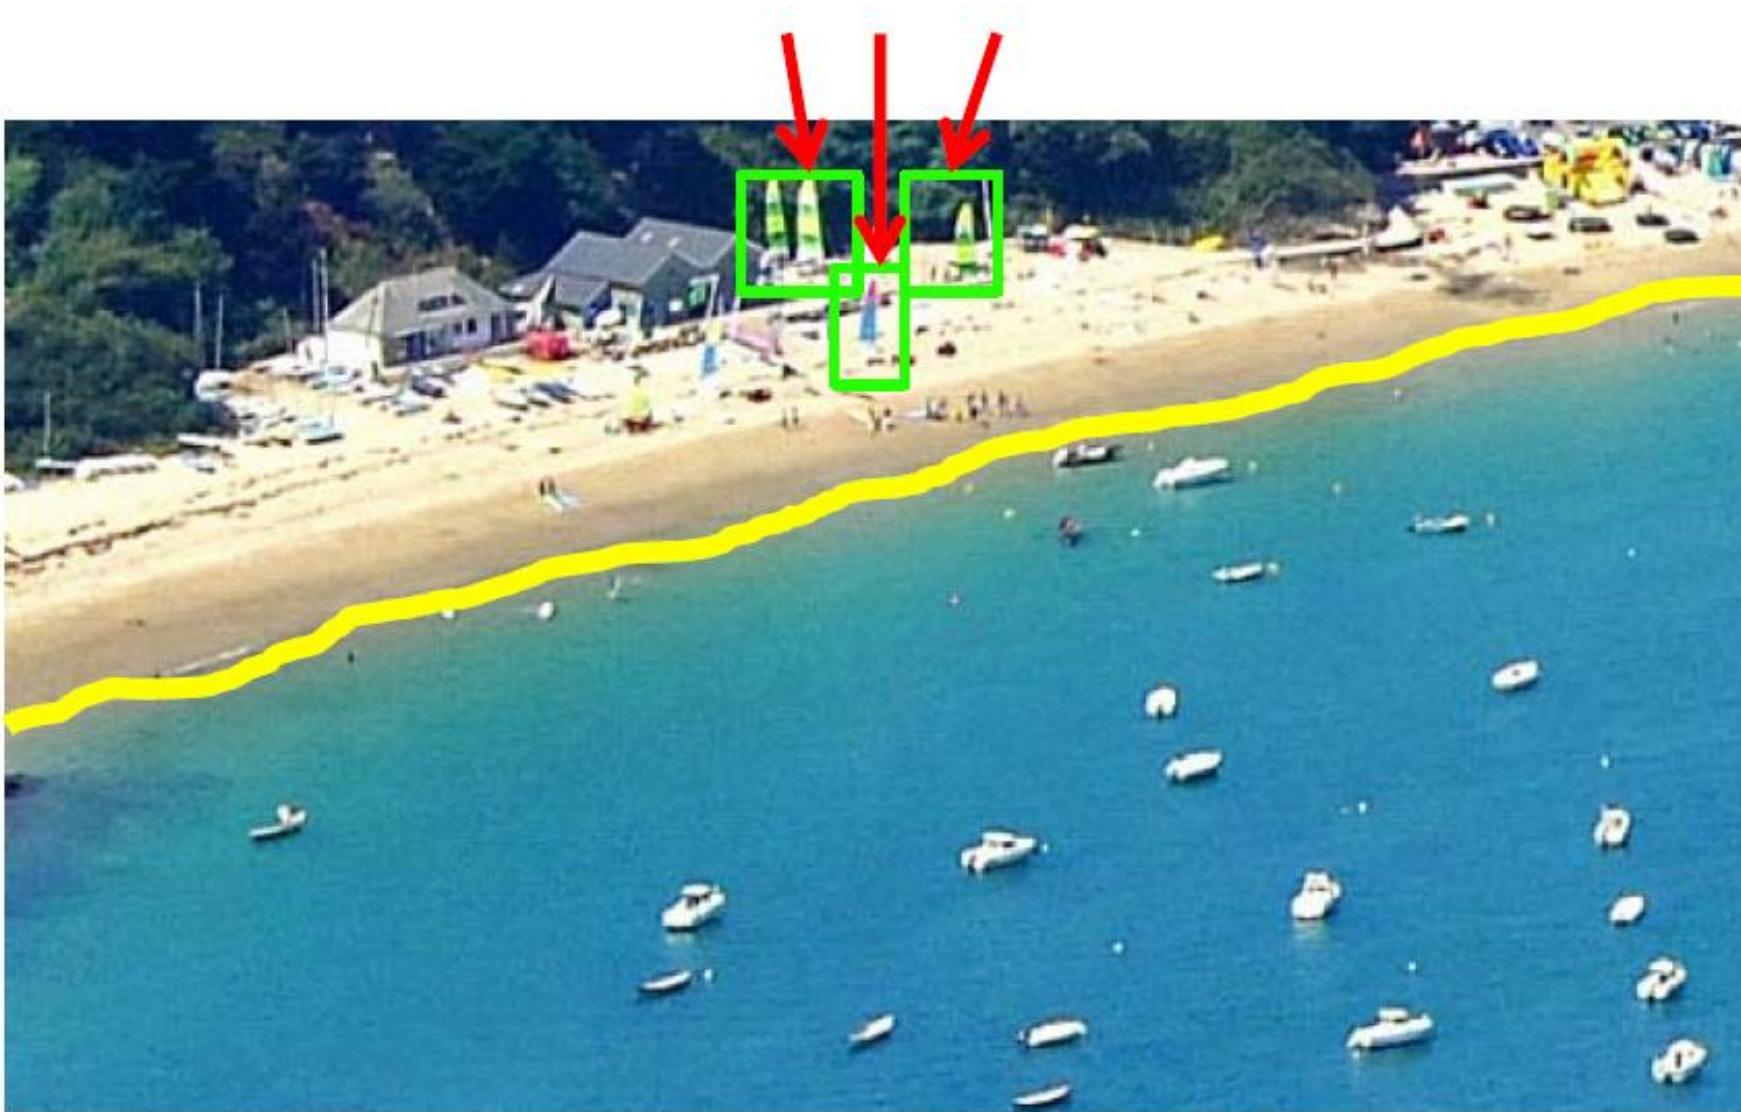
Classification: Basic Principles



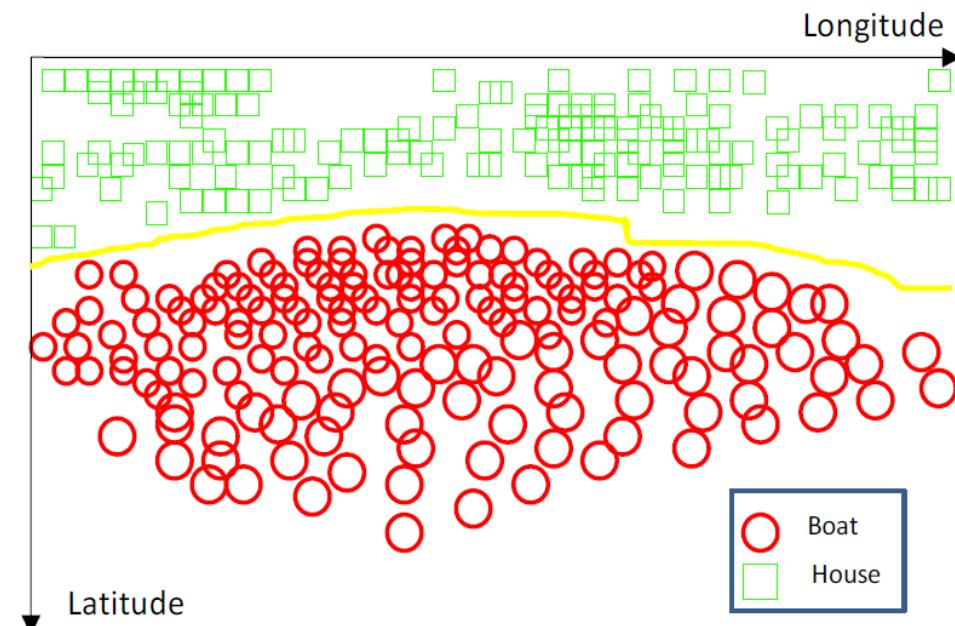
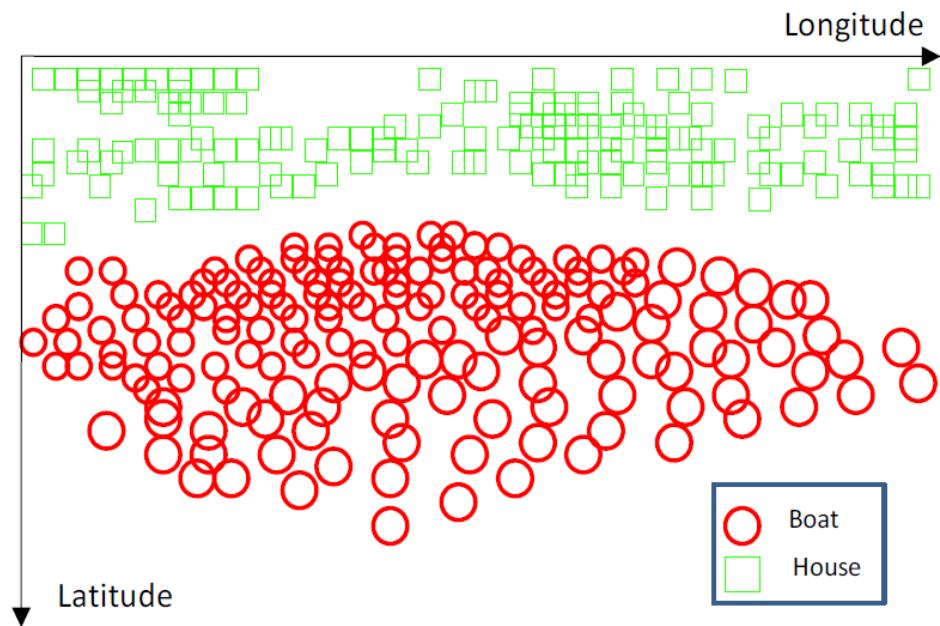
Classification: Decision Surface



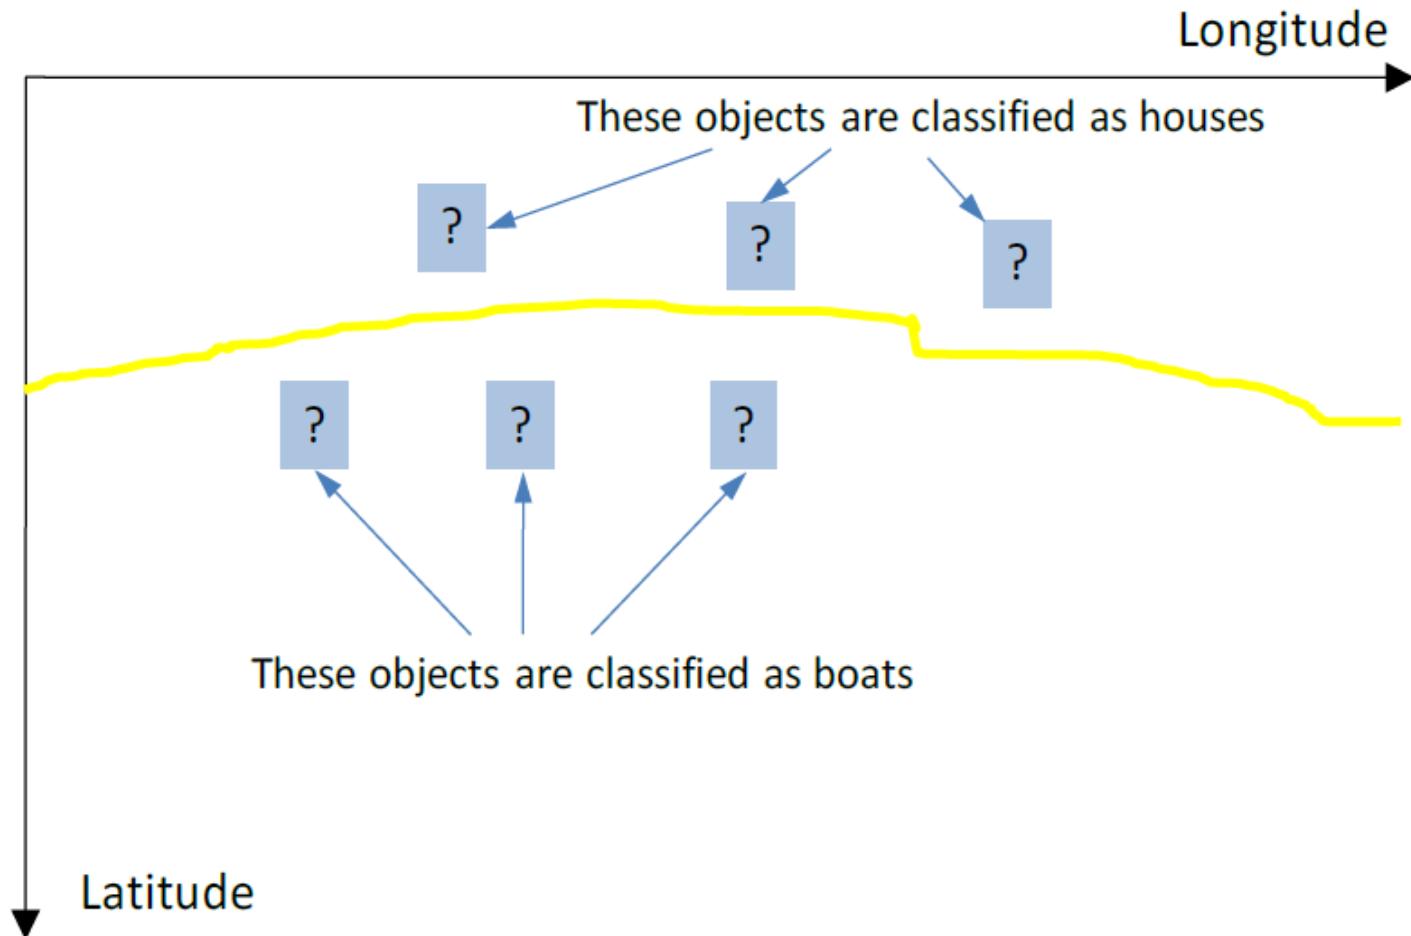
Classification: Classification Errors



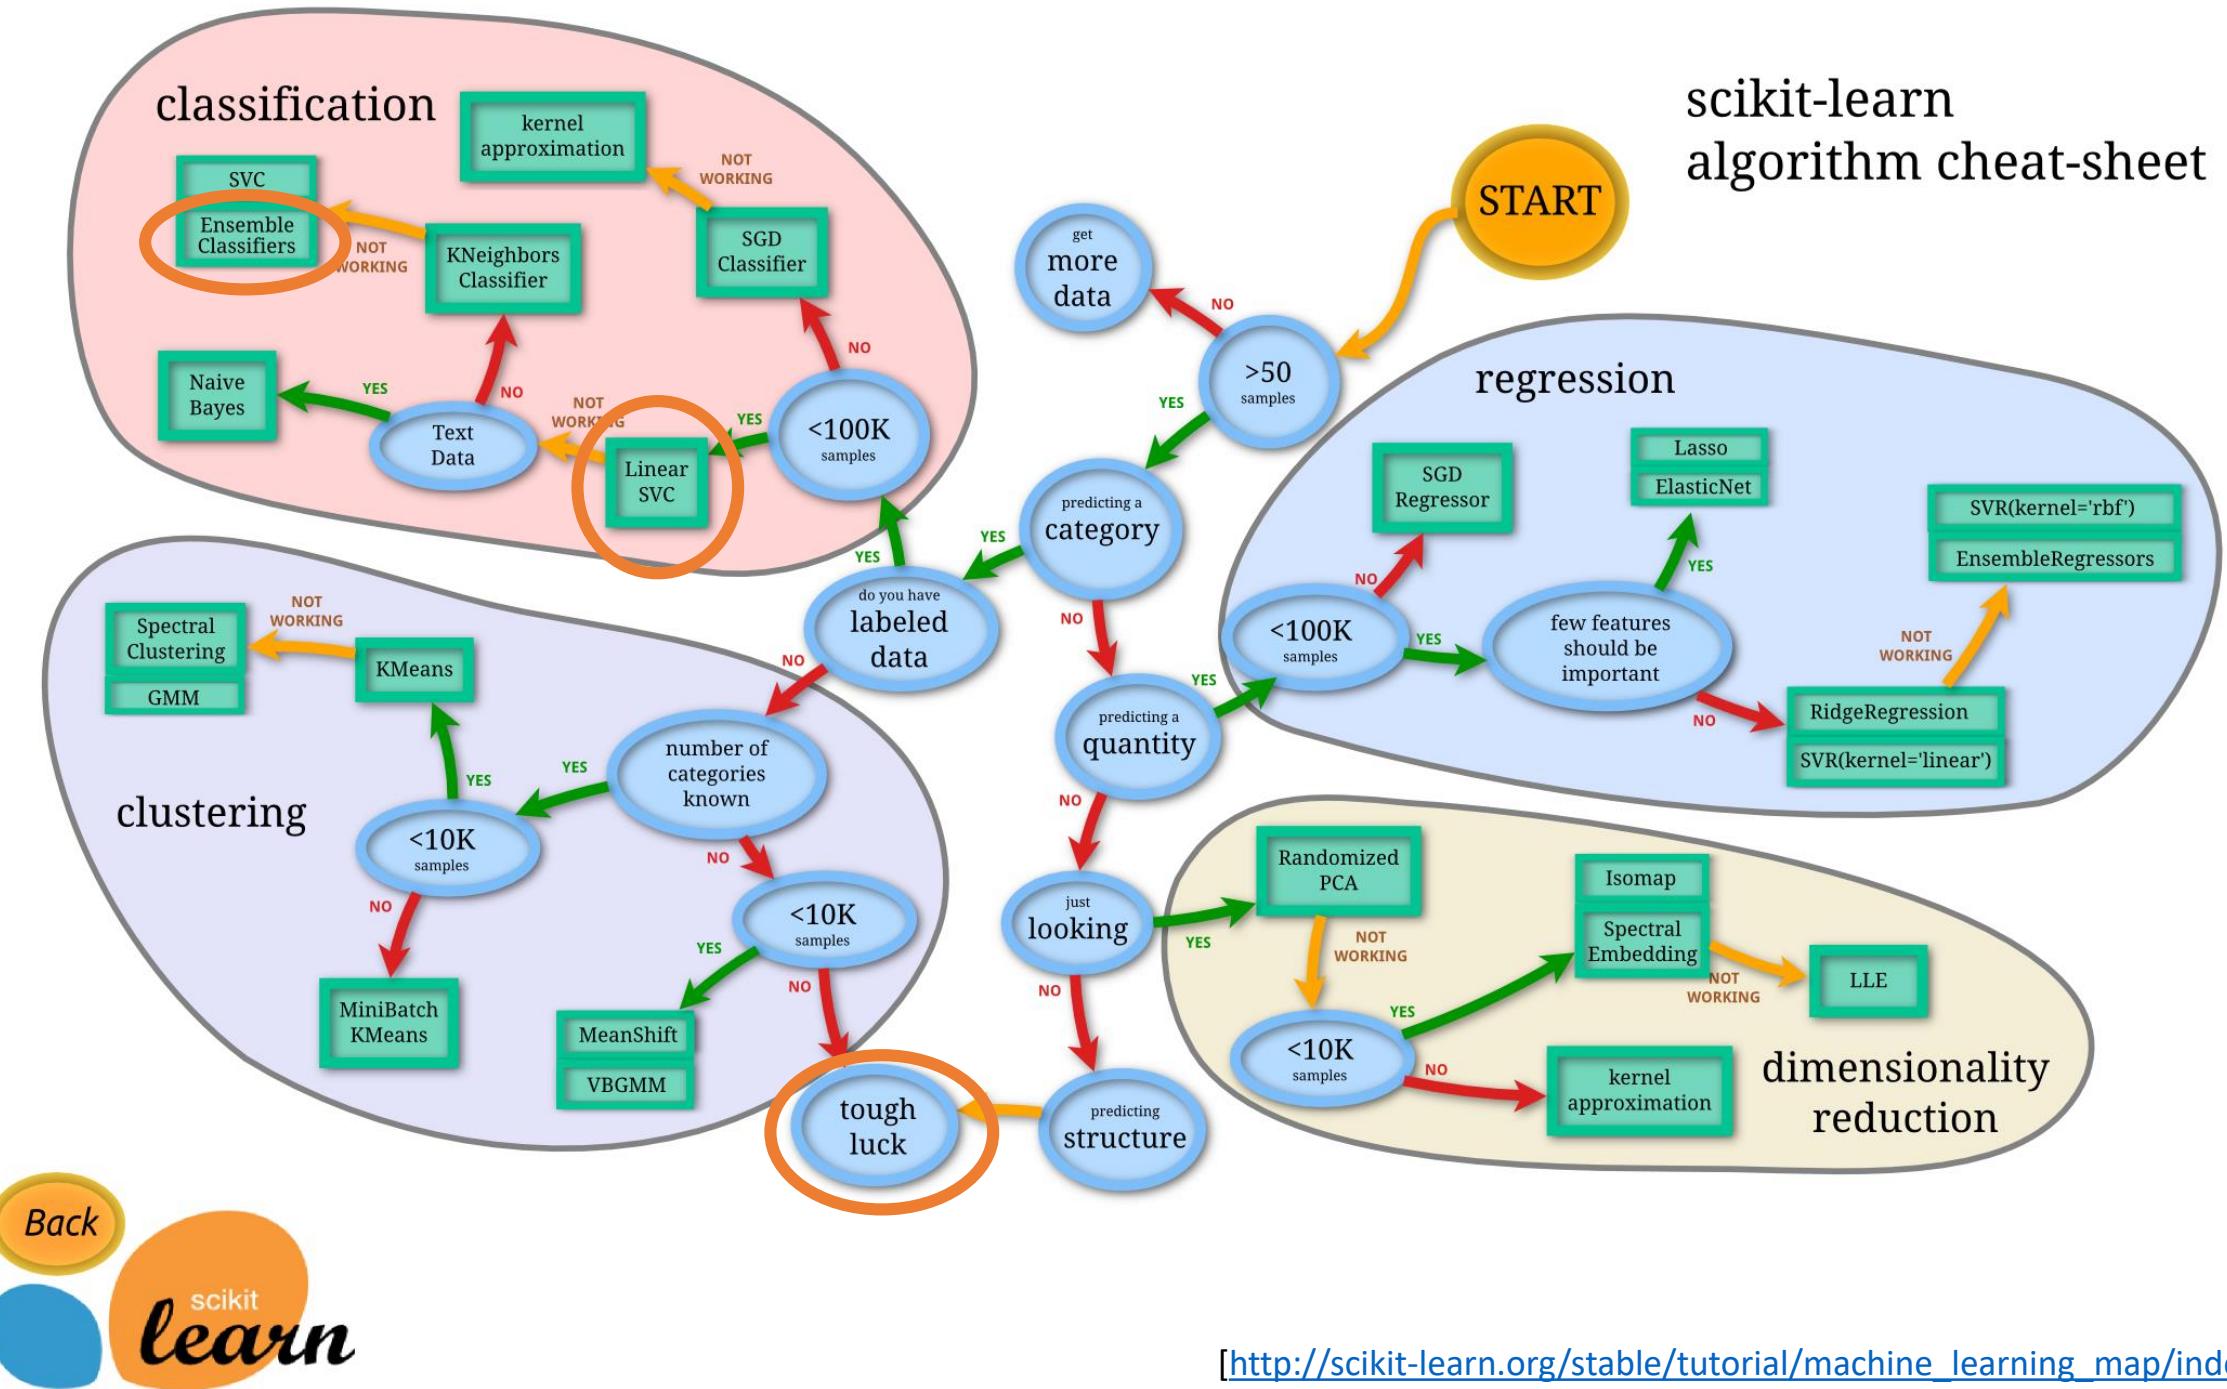
Classification Algorithms – Training time



Classification Algorithms – Test time



scikit-learn algorithm cheat-sheet



Support Vector Machine (SVM)

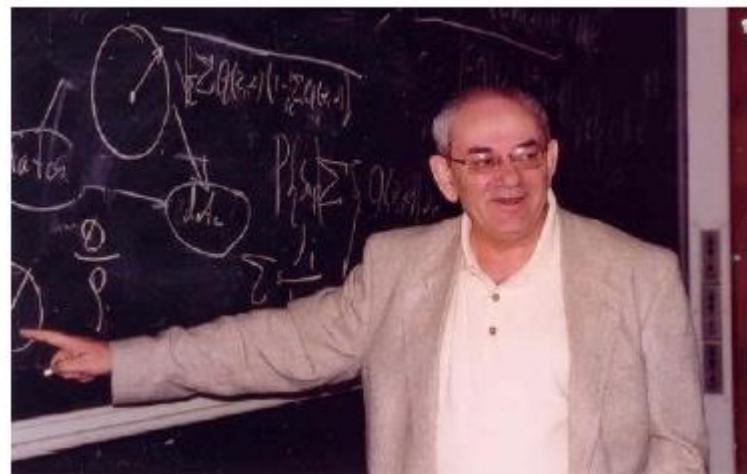
Derived from statistical learning theory by
Vapnik, et al. in 1992

Simple SVM achieved accuracy comparable to
neural-network with hand-designed features in
a handwriting recognition task

SVM widely used in object detection &
recognition, content-based image retrieval, text
recognition, biometrics, speech recognition, etc.

Also used for regression (will not cover today)

Vladimir Vapnik





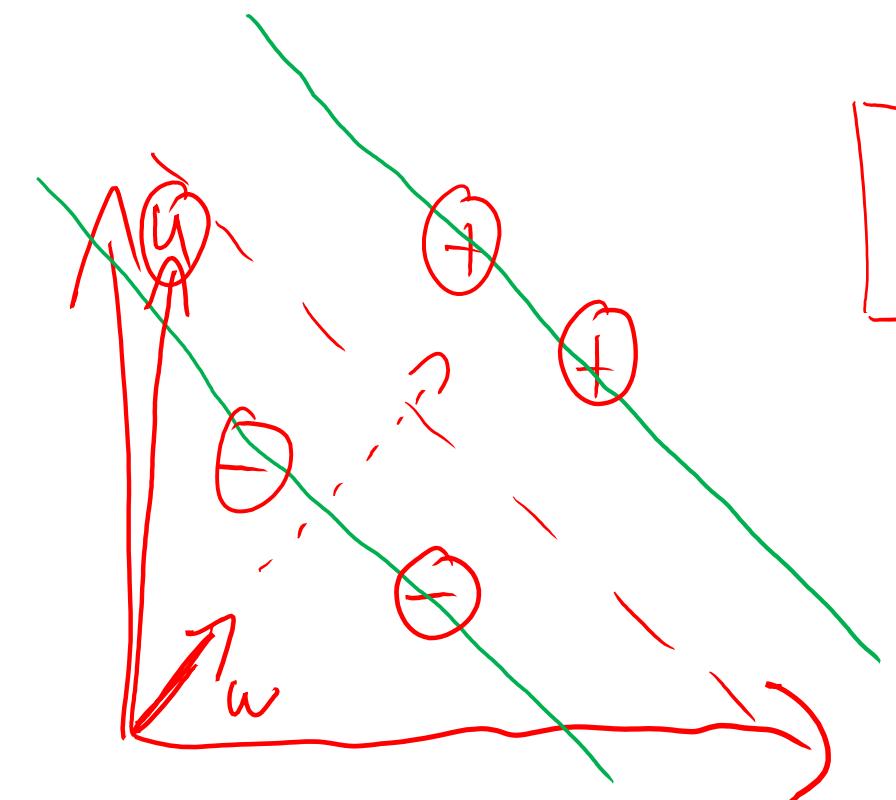
SVM Derivation (Intuition with some slight of hand)

$$\bar{w} \cdot \bar{u} \geq c$$

$$b = -c$$

$$\boxed{\bar{w} \cdot \bar{u} + b \geq 0}$$
 then \oplus

What are \bar{w} & b
are the unknowns



SVM Derivation (Intuition with some slight of hand)

$$\bar{\omega} \cdot \bar{x}_+ + b \geq 1$$

$$\bar{\omega} \cdot \bar{x}_- + b \leq -1$$

for convenience

we introduce

$$y \text{ such that } \begin{cases} +1 & \text{for } (+) \\ -1 & \text{for } (-) \end{cases}$$

$$y_i(\bar{\omega} \cdot \bar{x}_i + b) \geq 1 \text{ for SV}$$

$$\Rightarrow y_i(\bar{\omega} \cdot \bar{x}_i + b) \cancel{\geq 1} \Rightarrow y_i(\bar{\omega} \cdot \bar{x}_i + b) \geq 1$$

$$y_i(\bar{\omega} \cdot \bar{x}_i + b) = 1$$

$$y_i(\bar{\omega} \cdot \bar{x}_i +) - 1 = 0$$

SVM Derivation (Intuition with some slight of hand)

We need to constrain \bar{w} & b

$$\bar{w} \cdot \bar{x}_+ + b \stackrel{\text{positive sample}}{\geq} 1$$

$$\bar{w} \cdot \bar{x}_- + b \leq -1$$

↑
negative sample

$$y_i(\bar{w} \cdot \bar{x}_i + b) \geq 1$$

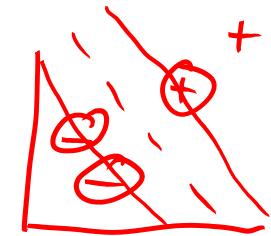
$$y_i(\bar{w} \cdot \bar{x}_i + b) \geq 1$$

$$y_i(\bar{w} \cdot \bar{x}_i + b) - 1 \geq 0$$

$$\boxed{y_i(\bar{w} \cdot \bar{x}_i + b) - 1 = 0}$$

FOR x_i ON THE CUPB

y_i such that
+ samples
- 1 for - samples



SVM Derivation (Intuition with some slight of hand)

$$y_i(\bar{w} \cdot \bar{x}_i + b) - 1 = 0$$



$$\text{width} = (\bar{x}_+ - \bar{x}_-) \cdot \frac{\bar{w}}{\|\bar{w}\|} = \frac{2}{\|\bar{w}\|} * \text{The primal problem}$$



$$x_+: g_i(\bar{w} \cdot \bar{x}_+ + b) = 0 \\ \bar{w} \cdot \bar{x}_+ = 1 - b$$

$$\max \frac{2}{\|\bar{w}\|} = \max \frac{1}{\|\bar{w}\|} = \min \|\bar{w}\| \approx \frac{1}{2} \|\bar{w}\|^2$$

SVM Derivation (Intuition with some slight of hand)

$$y_i(\bar{w} \cdot \bar{x}_i + b) - 1 = 0$$

$$\min \frac{1}{2} \|w\|^2$$

$$L = \frac{1}{2} \|\bar{w}\|^2 - \sum \alpha_i [y_i(\bar{w} \cdot \bar{x}_i + b) - 1]$$

↑ linear sum
↑ of samples \bar{x}_i

$$\frac{\partial L}{\partial w} = \bar{w} - \sum \alpha_i y_i x_i = 0 \Rightarrow \boxed{\bar{w} = \sum \alpha_i y_i \bar{x}_i}$$

$$\frac{\partial L}{\partial b} = -\sum \alpha_i y_i = 0 \Rightarrow \boxed{\sum \alpha_i y_i = 0}$$

$$L = \frac{1}{2} (\sum \alpha_i y_i \bar{x}_i)(\sum \alpha_j y_j \bar{x}_j) - (\sum \alpha_i y_i \bar{x}_i)(\sum \alpha_j y_j \bar{x}_j) - \underbrace{\sum \alpha_i y_i b}_{= 0} + \sum \alpha_i$$

$$= \sum \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j x_i \cdot x_j$$

SVM Derivation (Intuition with some slight of hand)

$$L = \sum \alpha_i - \frac{1}{2} \sum \sum \underbrace{\alpha_i \alpha_j}_{\text{Output}} \underbrace{y_i y_j}_{\text{coeffs from training data}} \boxed{\bar{x}_i \cdot \bar{x}_j}$$

optimization
only depends
on dot product
of sample
vecs

DECISION RULE:

$$\sum \alpha_i y_i \boxed{\bar{x}_i \cdot \bar{u}} + b \geq 0 \quad \text{THEN } +$$

decision also only depends on
dot product

The Optimization Problem Solution

The solution has the form:

$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i \quad b = y_k - \mathbf{w}^T \mathbf{x}_k \text{ for any } \mathbf{x}_k \text{ such that } \alpha_k \neq 0$$

Each non-zero α_i indicates that corresponding \mathbf{x}_i is a support vector.

Then the classifying function will have the form:

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

Optimization
is provably,
convex
(for linearly separable sets)

Notice that it relies only on the *dot product* between the test point \mathbf{x} and the support vectors \mathbf{x}_i .

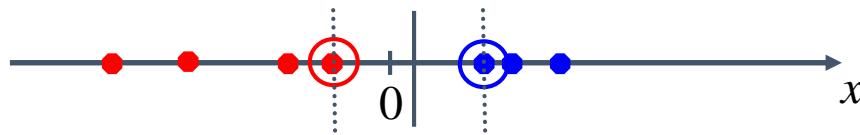
Also keep in mind that solving the optimization problem involved computing the inner products $\mathbf{x}_i^T \mathbf{x}_j$ between all pairs of training points.

Now Some Practical Issues (iPython Notebook)

Non-linear SVMs

Non-linear SVMs

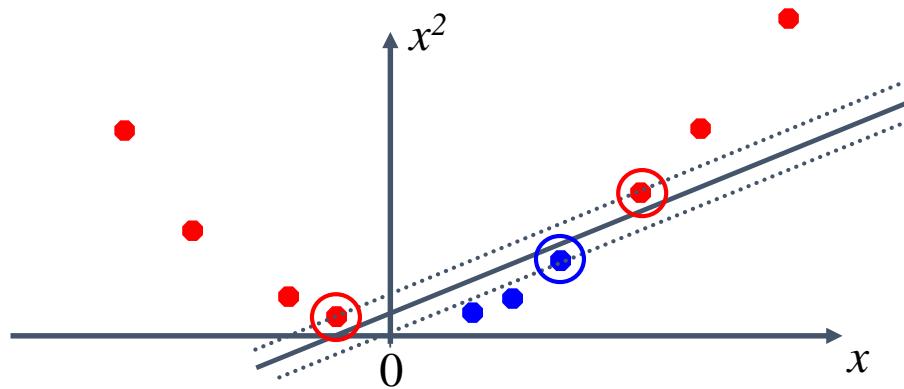
Datasets that are linearly separable (with some noise) work out great:



But what are we going to do if the dataset is just too hard?

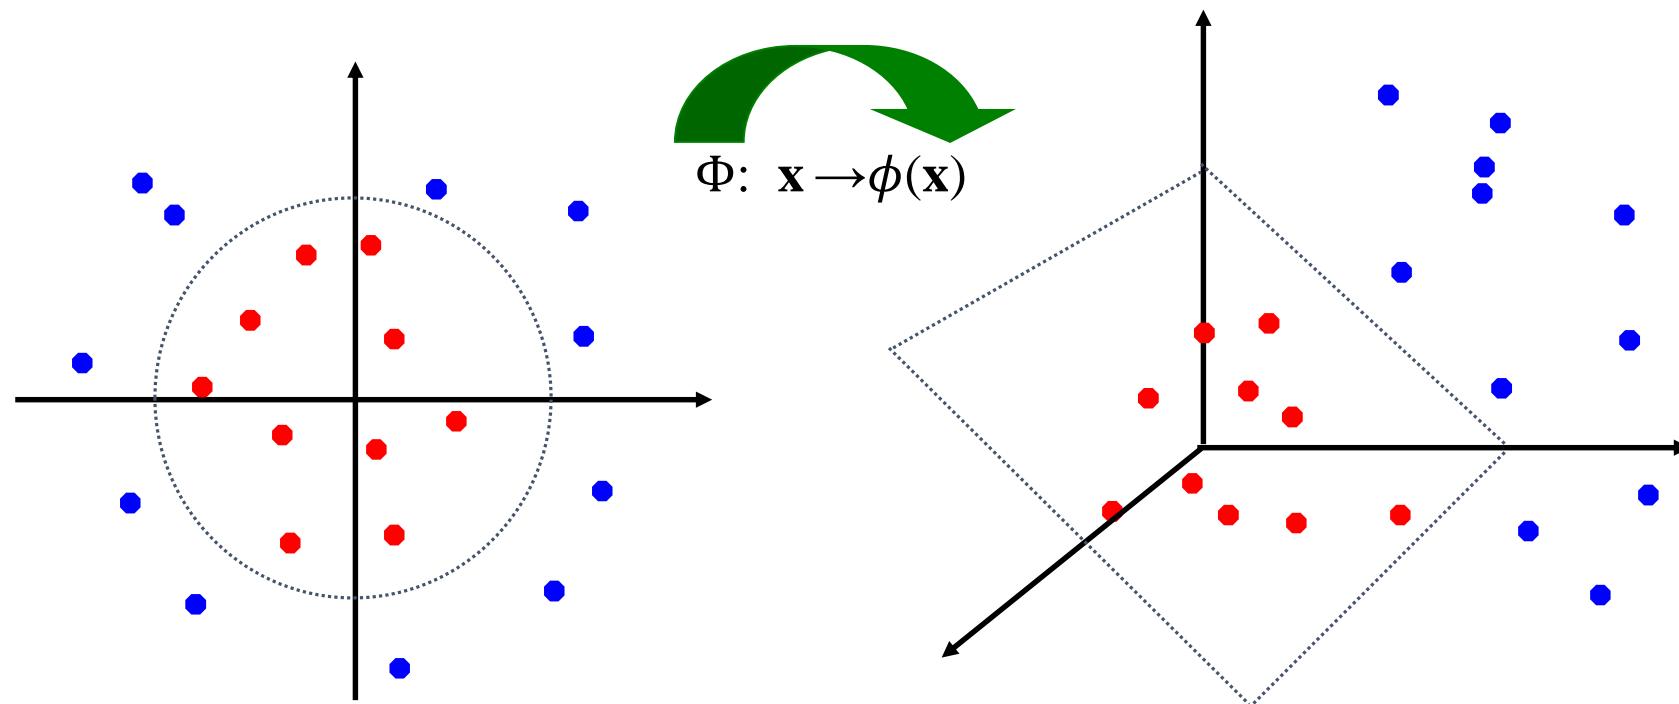


How about ... mapping data to a higher-dimensional space:



Non-linear SVMs: Feature spaces

General idea: the original feature space can always be mapped to some higher-dimensional feature space where the training set is linearly separable:



Non-linear SVMs: The Kernel Trick

With this mapping, our discriminant function is now:

$$y(x) = w^T \phi(x) + b = \sum_{i \in SV} \alpha_i \boxed{\phi(x_i)^T \phi(x)} + b$$

No need to know mapping function ϕ explicitly, because we only use the **dot product** of feature vectors in both training and test.

A *kernel function* is defined as a function that corresponds to a inner product of two feature vectors in some expanded feature space:

$$K(\mathbf{x}_i, \mathbf{x}_j) \equiv \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

Non-linear SVMs: The Kernel Trick

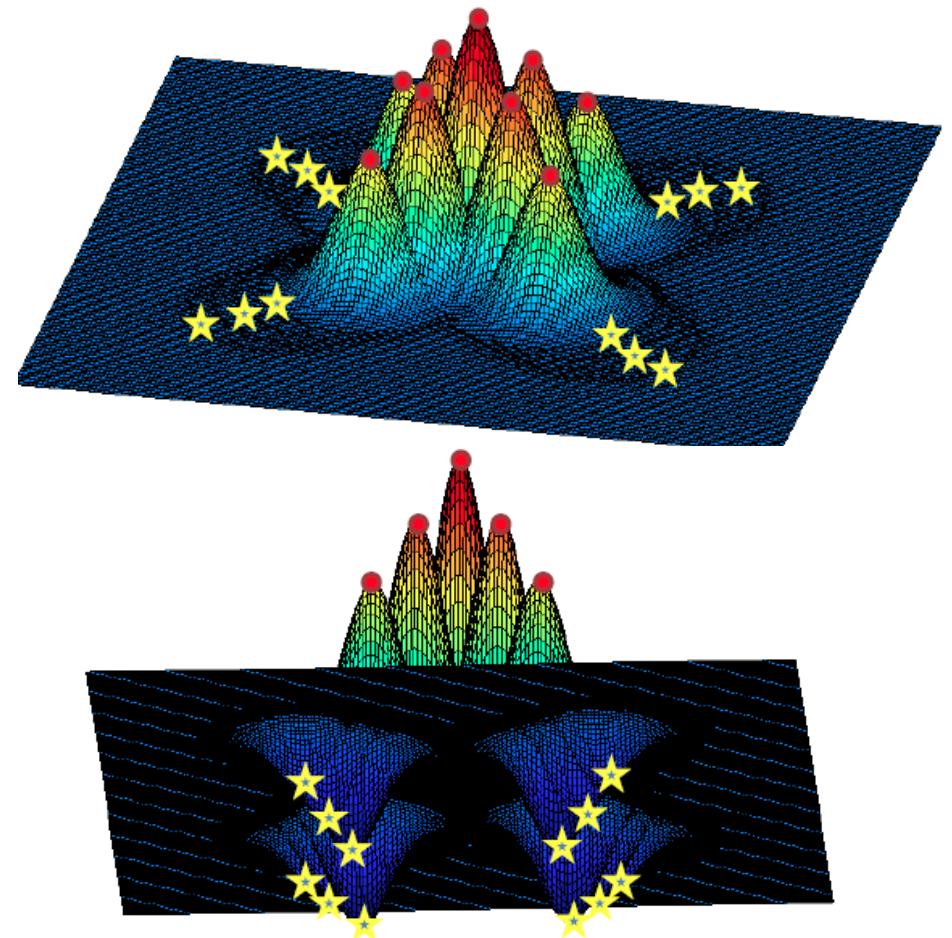
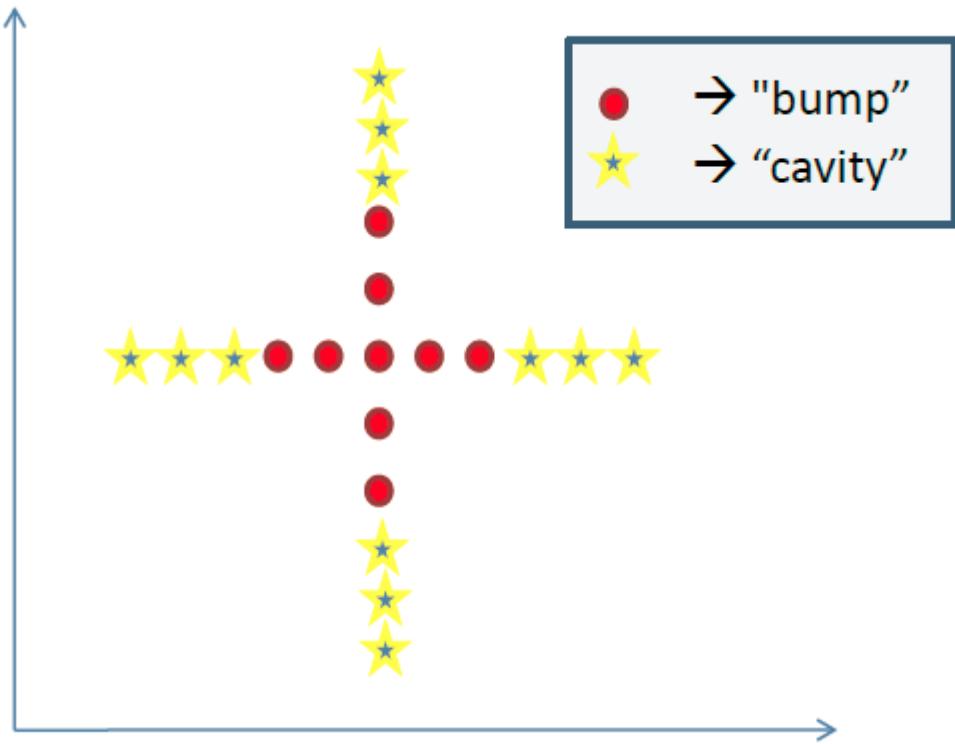
Commonly used kernel functions:

- Linear kernel: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- Polynomial kernel: $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^p$
- Gaussian (Radial-Basis Function Kernel (RBF)) kernel:
$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$
- Sigmoid: $K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\beta_0 \mathbf{x}_i^T + \beta_1)$

In general, functions that satisfy *Mercer's condition* can be kernel functions.

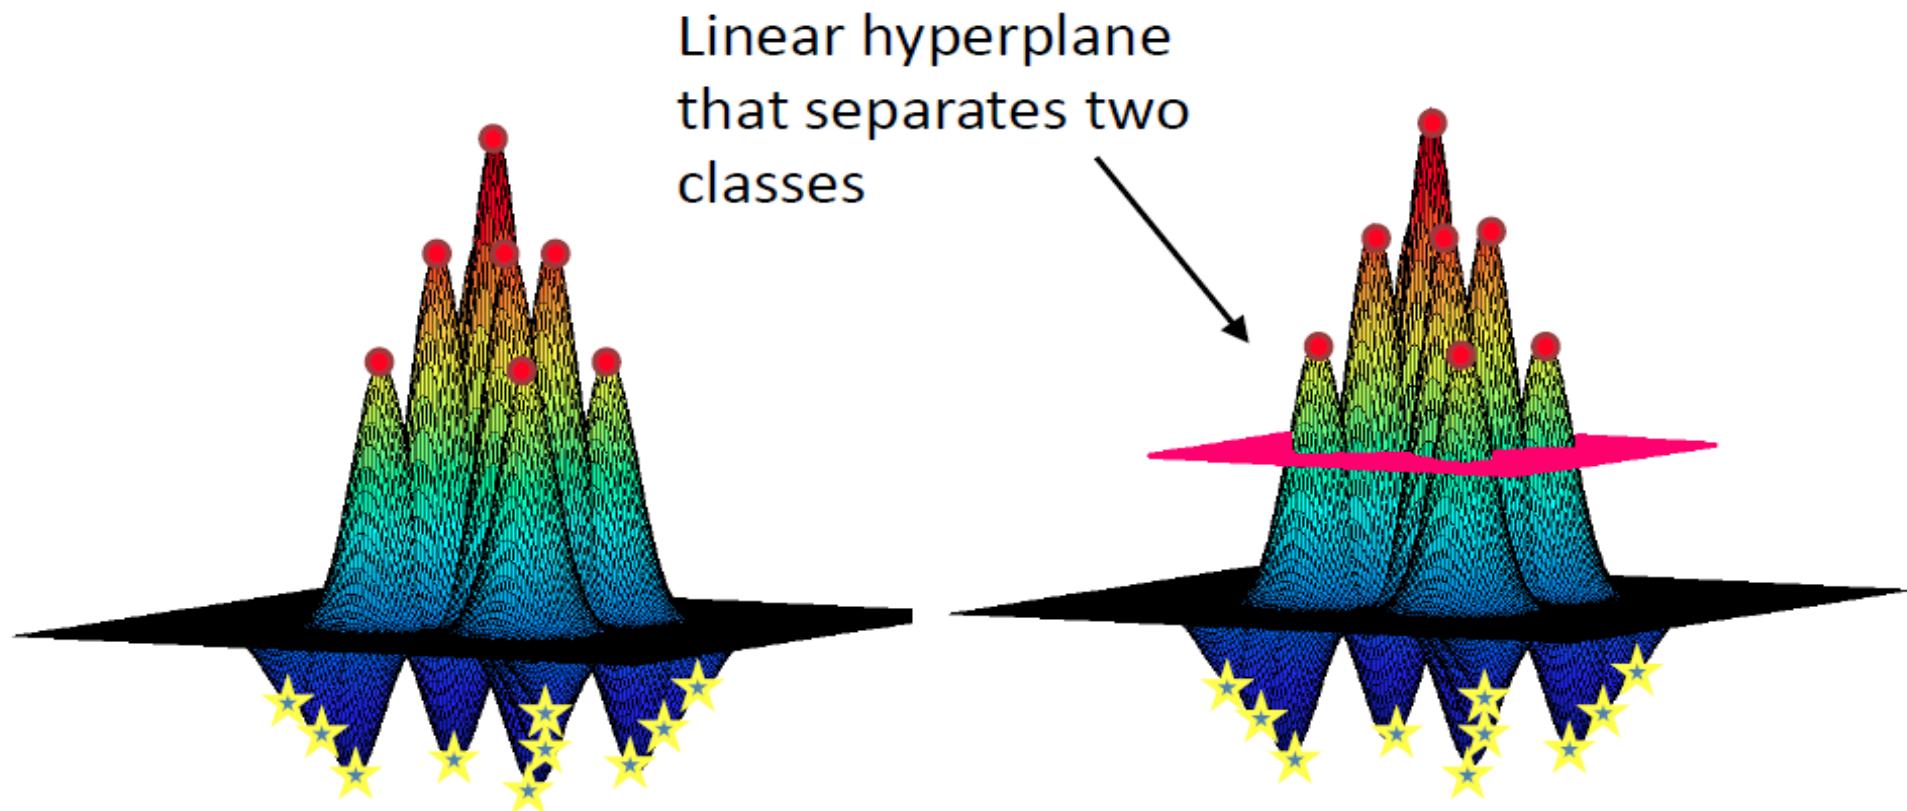
Example: Gaussian Kernel

Geometrically, this is a “bump” or “cavity” centered at the training data point x :



Example: Gaussian Kernel - Classification

Geometrically, this is a “bump” or “cavity” centered at the training data point x :



SVMs in HCI

Many software packages exist, implementing linear SVMs, kernel SVMs and more

- LibSVM (<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>)
- Matlab SVM toolbox
(<http://www.isis.ecs.soton.ac.uk/resources/svminfo/>)
- Weka Data Mining toolbox
(<http://www.cs.waikato.ac.nz/ml/weka/index.html>)
- More resources:
 - <http://www.kernel-machines.org>

SVMs probably *the most frequently* used ML technique in HCI

Reading Suggestions

V. Vapnik, Statistical Learning Theory. Wiley-Interscience; 1998

Christopher M. Bishop, Pattern Recognition and Machine Learning (Information Science and Statistics), Springer-Verlag New York, Inc., Secaucus, NJ, 2006

C.J.C. Burges. A tutorial on support vector machines for pattern recognition. Data Mining and Knowledge Discovery, 2(2):955-974, 1998.

<http://www.kernel-machines.org>